

3. Mirror cost under symmetric re-segmentation

OUTPUT *Abstract.* We define the *mirror cost* of an English span by removing its formatting, re-segmenting both its letter stream and its reversal with the same lexicon and procedure, and comparing causal-language-model cost per letter. This symmetric control isolates the effect of reversal from the effect of replacing author-supplied spaces. A positive gap means that, under the selected model and segmenter, the reversed reading is less likely. Exact decoding then illustrates the operational consequence of this gap.

CRITIQUE A likelihood gap is not a theorem about English. It does not count grammatical strings, show that a palindrome is impossible, measure meaning, or justify information-theoretic language.

REVISION APPLIED State the metric narrowly as a model-based diagnostic, and remove claims about feasibility counts, grammaticality, human preference, and lower bounds.