

5. A bounded empirical claim about reversed English

OUTPUT *Abstract.* On 57 retained WikiText-2 configurations, reverse readings incur 2.16–3.60 additional bits per letter after matched re-segmentation. The experiment varies span length, three segmentation objectives, and DistilGPT-2, GPT-2, and GPT-2-medium, with a limited GPT-2-large check. The results indicate a stable reverse-reading penalty within this design; they do not establish that every language model or every English corpus behaves identically.

CRITIQUE These checkpoints are all GPT-2-family models. Reporting them as independent architecture replications would overstate the evidence. A single corpus and lexicon are similarly too narrow for a universal English claim.

REVISION APPLIED Describe model variation as a scale-sensitivity check, name the corpus and lexicon in the claim, and make cross-corpus and cross-architecture replication a limitation.